

STAT619 project Analysis 2 - Nasdaq-100/QQQ Pricing

Erik Duus

5/5/2021

Retrieve data from Yahoo finance:

- use simple interpolation to approximate missing values.
- save as a local file so download not required each time

```
options("getSymbols.warning4.0"=FALSE)
options("getSymbols.yahoo.warning"=FALSE)

# get data from YHOO and save to disk

qqq <- getSymbols("QQQ", auto.assign = FALSE)

# which(is.na(btc$'BTC-USD.Close'))
qqq <- na.approx(qqq)


#saveRDS(qqq, './qqq.rds')
```

Load local copy of pricing data

- Examine time series; it is growing exponentially
- Log transform and difference to convert to a returns series
- Mean looks stationary, variance is not constant
- Returns show volatility - GARCH may be needed
- Extract the 2016-2017 window for analysis
- Plot ACF/PACF of returns and squared returns; squared returns confirm GARCH

```
# load saved Yahoo data; use adjusted closing prices
qqqall = readRDS('./qqq.rds')
qqqall <- qqqall$QQQ.Adjusted


qqqr <- diff(log(qqqall))
plot(qqqr, main='QQQ daily returns')
```



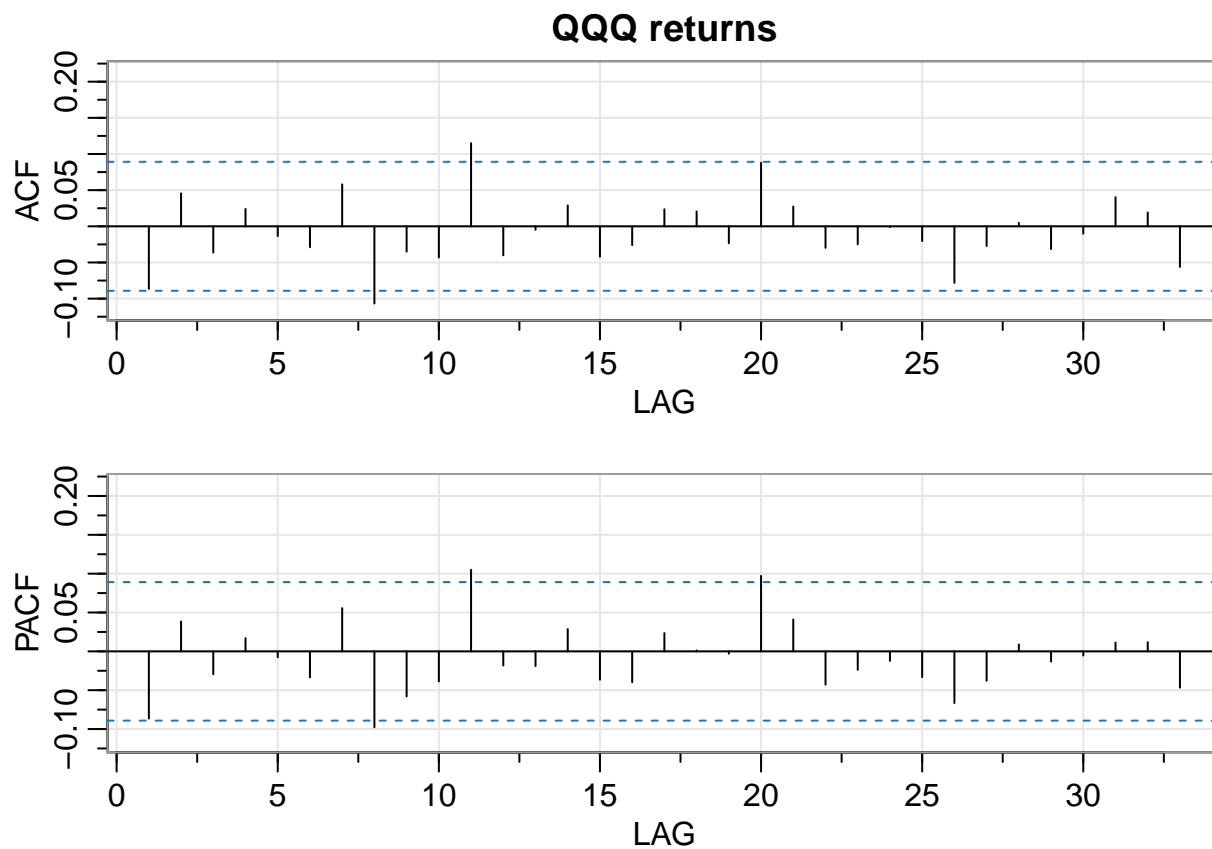
```
qqq <- window(qqqall, start='2016-01-01',end='2018-01-01')
qqqr <- window(qqqr, start='2016-01-01',end='2018-01-01')

layout(matrix(c(1,1,1,1,
                2,2,2,2),
              nrow=2, byrow=TRUE))

plot(qqqall, main='QQQ closing price')
plot(qqqr, main='QQQ daily returns 2016-2017')
```

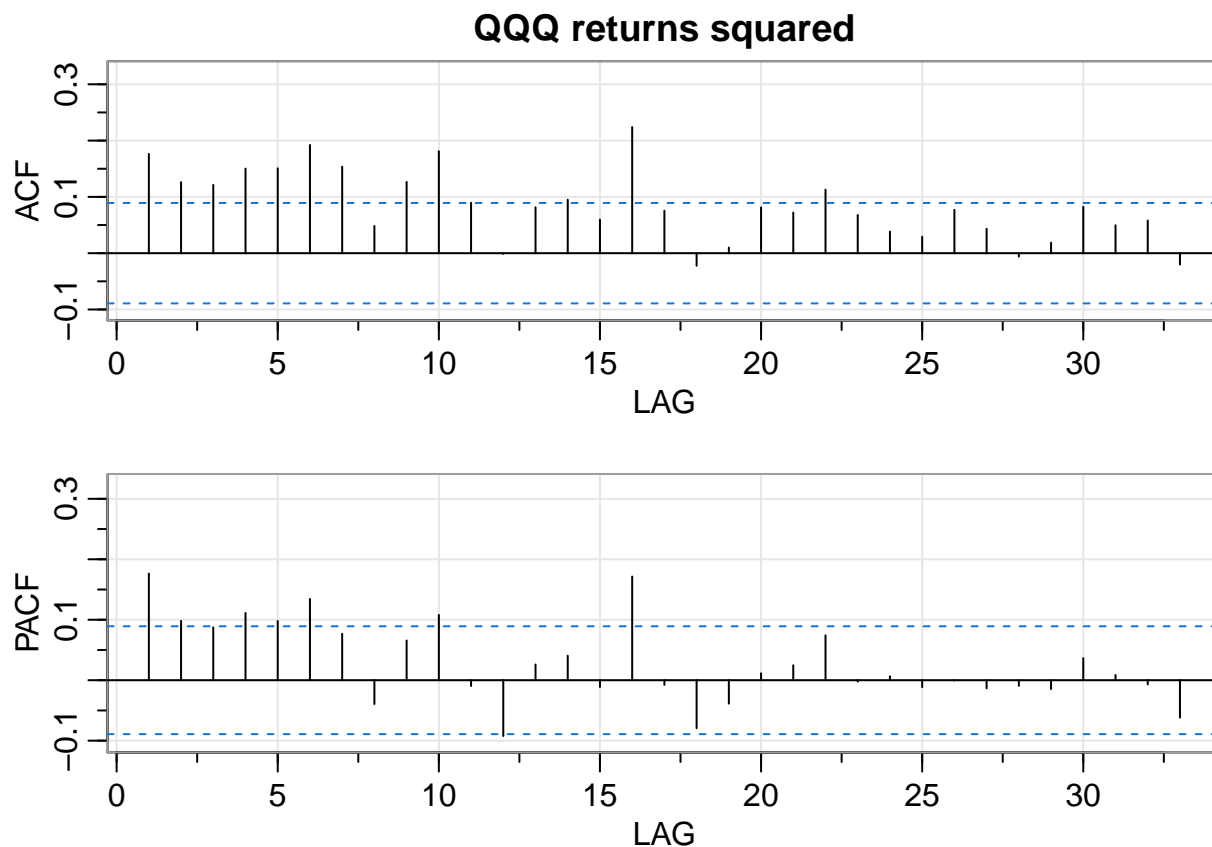


```
acf2(qqqr, main='QQQ returns')
```



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
## ACF  -0.09 0.05 -0.04 0.02 -0.01 -0.03 0.06 -0.11 -0.03 -0.04 0.12 -0.04 0.00
## PACF  -0.09 0.04 -0.03 0.02 -0.01 -0.03 0.06 -0.10 -0.06 -0.04 0.11 -0.02 -0.02
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]
## ACF   0.03 -0.04 -0.03 0.02 0.02 -0.02 0.09 0.03 -0.03 -0.02 0.00 -0.02
## PACF   0.03 -0.04 -0.04 0.02 0.00 0.00 0.10 0.04 -0.04 -0.02 -0.01 -0.03
##      [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33]
## ACF  -0.08 -0.03 0.00 -0.03 -0.01 0.04 0.02 -0.06
## PACF -0.07 -0.04 0.01 -0.01 -0.01 0.01 0.01 -0.05
```

```
acf2(qqqr2, main='QQQ returns squared')
```



```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7]  [,8] [,9] [,10] [,11] [,12] [,13]
## ACF  0.18 0.13 0.12 0.15 0.15 0.19 0.15  0.05 0.13  0.18 0.09 0.00 0.08
## PACF 0.18 0.10 0.09 0.11 0.10 0.13 0.08 -0.04 0.07  0.11 -0.01 -0.09 0.03
##      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]
## ACF  0.10 0.06 0.22 0.08 -0.02 0.01 0.08 0.07 0.11 0.07 0.04 0.03
## PACF 0.04 -0.01 0.17 -0.01 -0.08 -0.04 0.01 0.02 0.07 0.00 0.01 -0.01
##      [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33]
## ACF  0.08 0.04 -0.01 0.02 0.08 0.05 0.06 -0.02
## PACF 0.00 -0.01 -0.01 -0.01 0.04 0.01 -0.01 -0.06
```

confirm stationarity of mean through unit root tests

```
# null hypothesis not stationary
adf.test(qqqr, k=0)
```

```
## Warning in adf.test(qqqr, k = 0): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: qqqr
## Dickey-Fuller = -24.551, Lag order = 0, p-value = 0.01
## alternative hypothesis: stationary
```

```
adf.test(qqqr)
```

```
## Warning in adf.test(qqqr): p-value smaller than printed p-value
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: qqqr  
## Dickey-Fuller = -9.3445, Lag order = 7, p-value = 0.01  
## alternative hypothesis: stationary
```

```
pp.test(qqqr)
```

```
## Warning in pp.test(qqqr): p-value smaller than printed p-value
```

```
##  
## Phillips-Perron Unit Root Test  
##  
## data: qqqr  
## Dickey-Fuller Z(alpha) = -543.77, Truncation lag parameter = 5, p-value  
## = 0.01  
## alternative hypothesis: stationary
```

```
# null hypothesis stationary
```

```
kpss.test(qqqr)
```

```
## Warning in kpss.test(qqqr): p-value greater than printed p-value
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: qqqr  
## KPSS Level = 0.14991, Truncation lag parameter = 5, p-value = 0.1
```


Run `auto.arima` to see what it suggests;diagnostics of fitted model also show variance in residuals, confirming garch.

```
auto.arima(qqqr, seasonal=FALSE)
```

```
## Series: qqqr  
## ARIMA(1,0,0) with non-zero mean  
##  
## Coefficients:  
##          ar1    mean  
##      -0.0876 7e-04  
## s.e.   0.0447 4e-04  
##  
## sigma^2 estimated as 7.302e-05: log likelihood=1682.76  
## AIC=-3359.52 AICc=-3359.47 BIC=-3346.86
```

```
sarima(qqqr, 1,0,0)
```

```
## initial value -4.766209
## iter 2 value -4.770027
## iter 3 value -4.770031
## iter 4 value -4.770036
## iter 5 value -4.770037
## iter 5 value -4.770037
## iter 5 value -4.770037
## final value -4.770037
## converged
## initial value -4.764377
## iter 2 value -4.764382
## iter 3 value -4.764388
## iter 4 value -4.764388
## iter 4 value -4.764388
## iter 4 value -4.764388
## final value -4.764388
## converged
```



```
## $fit
##
## Call:
## arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D, Q), period = S),
```

```
##      xreg = xmean, include.mean = FALSE, transform.pars = trans, fixed = fixed,
##      optim.control = list(trace = trc, REPORT = 1, reltol = tol))
##
## Coefficients:
##          ar1  xmean
##      -0.0876  7e-04
## s.e.    0.0447  4e-04
##
## sigma^2 estimated as 7.273e-05:  log likelihood = 1682.76,  aic = -3359.52
##
## $degrees_of_freedom
## [1] 501
##
## $ttable
##      Estimate      SE t.value p.value
## ar1    -0.0876 0.0447 -1.9602  0.0505
## xmean    0.0007 0.0004  1.9938  0.0467
##
## $AIC
## [1] -6.678971
##
## $AICc
## [1] -6.678924
##
## $BIC
## [1] -6.653799
```

ACF and PACF of series shows a tiny bit of auto-correlation, suggesting ARMA(1,1)

ACF and PACF of squared series shows both decaying, suggesting GARCH(1,1)

Fit GARCH(1,1) to start as a baseline

```
gf <- garchFit(~garch(1,1), data=qqqr, cond.dist='std', trace=FALSE)
```

```
## Warning: Using formula(x) is deprecated when x is a character vector of length > 1.
##   Consider formula(paste(x, collapse = " ")) instead.
```

```
summary(gf)
```

```
##
## Title:
##   GARCH Modelling
##
## Call:
##   garchFit(formula = ~garch(1, 1), data = qqqr, cond.dist = "std",
##     trace = FALSE)
##
## Mean and Variance Equation:
##   data ~ garch(1, 1)
## <environment: 0x0000000021ff0328>
##   [data = qqqr]
##
```



```

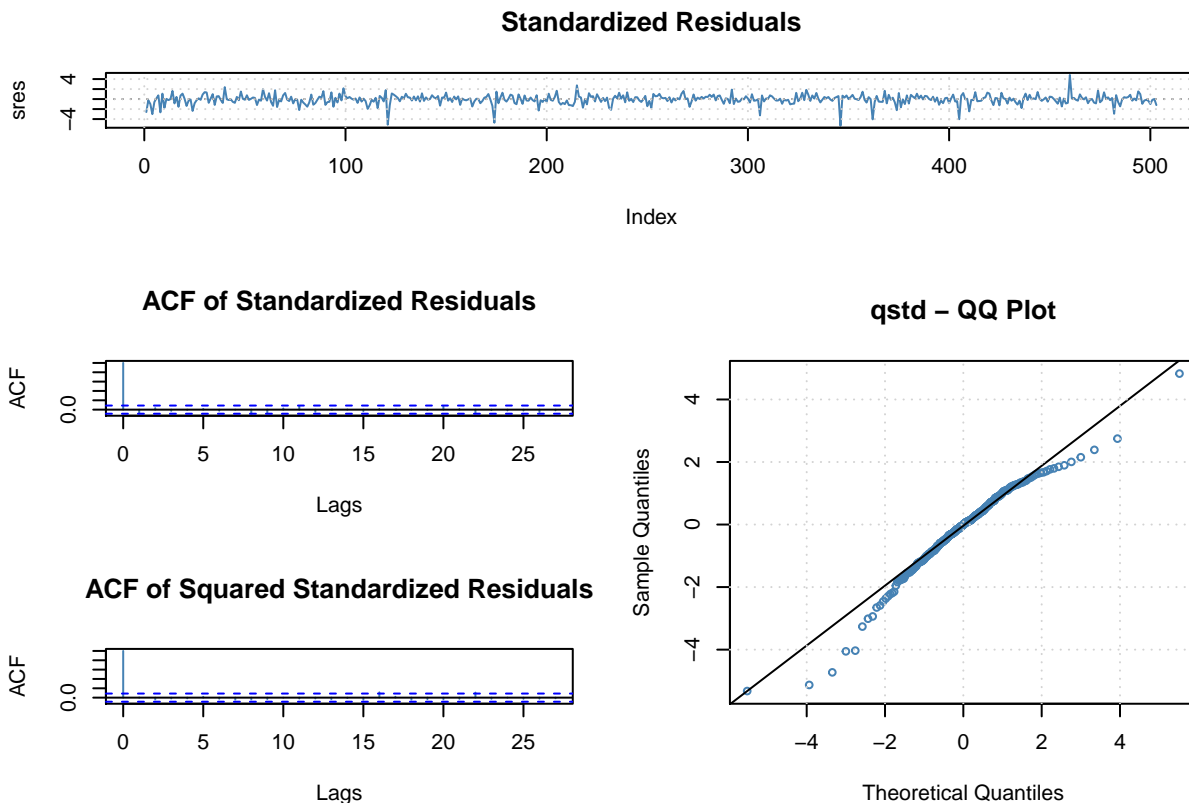
## Conditional Distribution:
## std
##
## Coefficient(s):
##      mu      omega      alpha1      beta1      shape
## 1.1094e-03 2.5634e-06 1.3668e-01 8.5113e-01 3.4234e+00
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      1.109e-03 2.560e-04 4.334 1.46e-05 ***
## omega  2.563e-06 1.683e-06 1.523 0.12767
## alpha1 1.367e-01 5.250e-02 2.604 0.00923 **
## beta1  8.511e-01 4.806e-02 17.709 < 2e-16 ***
## shape  3.423e+00 5.608e-01 6.104 1.03e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## 1763.356      normalized: 3.505678
##
## Description:
## Tue May 25 19:42:04 2021 by user: orovi
##
## Standardised Residuals Tests:
##
##      Statistic p-Value
## Jarque-Bera Test  R      Chi^2 718.8286 0
## Shapiro-Wilk Test R      W      0.9216501 1.64025e-15
## Ljung-Box Test    R      Q(10) 10.40282 0.4058908
## Ljung-Box Test    R      Q(15) 15.14452 0.4410579
## Ljung-Box Test    R      Q(20) 17.57307 0.6155093
## Ljung-Box Test    R^2 Q(10) 3.521033 0.9663794
## Ljung-Box Test    R^2 Q(15) 5.491048 0.9870892
## Ljung-Box Test    R^2 Q(20) 14.55135 0.8014781
## LM Arch Test      R      TR^2 5.347455 0.945356
##
## Information Criterion Statistics:
##      AIC      BIC      SIC      HQIC
## -6.991476 -6.949522 -6.991671 -6.975018

```

```

layout(matrix(c(1,1,1,1,
                1,1,1,1,
                2,2,4,4,
                2,2,4,4,
                3,3,4,4,
                3,3,4,4),nrow=6, byrow=TRUE))
plot(gf, which=9)
plot(gf, which=10)
plot(gf, which=11)
plot(gf, which=13)

```



model coefficients are significant, but QQ-plot is not that normal. perhaps ARMA features are required.
several iterations get us to ARMA(5,3)+GARCH(1,1)

**** (NOTE ARMA(5,3) is best model but wont predict correctly) ****

```
gf53 <- garchFit(~arma(5,3)+garch(1,1), data=qqqr ,cond.dist='std', trace=FALSE)
```

```
## Warning: Using formula(x) is deprecated when x is a character vector of length > 1.  
## Consider formula(paste(x, collapse = " ")) instead.
```

```
summary(gf53)
```

```
##  
## Title:  
## GARCH Modelling  
##  
## Call:  
## garchFit(formula = ~arma(5, 3) + garch(1, 1), data = qqqr, cond.dist = "std",  
## trace = FALSE)  
##  
## Mean and Variance Equation:  
## data ~ arma(5, 3) + garch(1, 1)  
## <environment: 0x00000000203f5f40>  
## [data = qqqr]
```

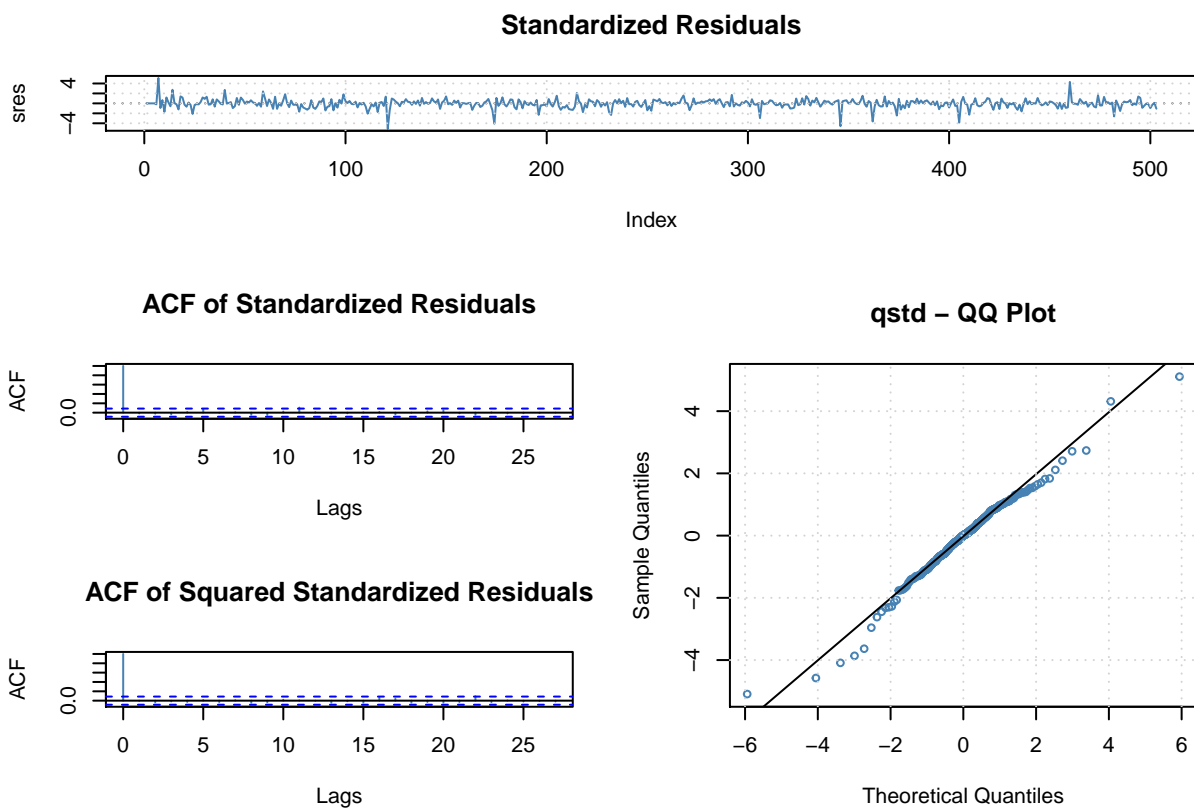
```

##
## Conditional Distribution:
## std
##
## Coefficient(s):
##      mu      ar1      ar2      ar3      ar4      ar5
## 2.6551e-04  4.2359e-01 -3.9607e-01  7.7965e-01  9.3830e-02 -1.2449e-01
##      ma1      ma2      ma3      omega      alpha1      beta1
## -5.1537e-01  5.2407e-01 -8.9357e-01  4.4330e-06  1.9124e-01  8.1233e-01
##      shape
## 2.9658e+00
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      2.655e-04  7.909e-05  3.357 0.000788 ***
## ar1      4.236e-01  5.116e-02  8.280 2.22e-16 ***
## ar2     -3.961e-01  5.100e-02 -7.767 7.99e-15 ***
## ar3      7.796e-01  4.196e-02 18.581 < 2e-16 ***
## ar4      9.383e-02  4.548e-02  2.063 0.039090 *
## ar5     -1.245e-01  4.165e-02 -2.989 0.002797 **
## ma1     -5.154e-01  2.697e-02 -19.112 < 2e-16 ***
## ma2      5.241e-01  2.898e-02 18.086 < 2e-16 ***
## ma3     -8.936e-01  2.901e-02 -30.804 < 2e-16 ***
## omega    4.433e-06  2.821e-06  1.572 0.116065
## alpha1   1.912e-01  8.732e-02  2.190 0.028522 *
## beta1    8.123e-01  6.315e-02 12.863 < 2e-16 ***
## shape    2.966e+00  5.067e-01  5.854 4.81e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## 1783.878      normalized: 3.546477
##
## Description:
## Tue May 25 19:42:05 2021 by user: orovi
##
##
## Standardised Residuals Tests:
##
##      Statistic p-Value
## Jarque-Bera Test  R      Chi^2  993.4853  0
## Shapiro-Wilk Test  R      W      0.9085347  0
## Ljung-Box Test     R      Q(10)  6.499543  0.7716946
## Ljung-Box Test     R      Q(15) 12.9058  0.6095736
## Ljung-Box Test     R      Q(20) 15.32202  0.7576944
## Ljung-Box Test     R^2    Q(10)  3.644229  0.9619728
## Ljung-Box Test     R^2    Q(15)  5.527343  0.9866417
## Ljung-Box Test     R^2    Q(20) 10.38313  0.9606809
## LM Arch Test       R      TR^2   4.768475  0.9652682
##
## Information Criterion Statistics:
##      AIC      BIC      SIC      HQIC

```

```
## -7.041263 -6.932182 -7.042555 -6.998471
```

```
layout(matrix(c(1,1,1,1,
                1,1,1,1,
                2,2,4,4,
                2,2,4,4,
                3,3,4,4,
                3,3,4,4),nrow=6, byrow=TRUE))
plot(gf53, which=9)
plot(gf53, which=10)
plot(gf53, which=11)
plot(gf53, which=13)
```



QQ plot is much better. residual plot and ACF plots suggest residuals are white noise. Ljung-Box tests are also suggestive of white noise residuals.

Conclude this model is accurate and use for predictions.

NOTE: fgarch package produces errors on this model when forecasting! revert to ARMA(2,2)+GARCH(1,1) for forecast

```
gf1 <- garchFit(~arma(2,2)+garch(1,1), data=qqqr ,cond.dist='std', trace=FALSE)
```

```
## Warning in log(s2): NaNs produced
```

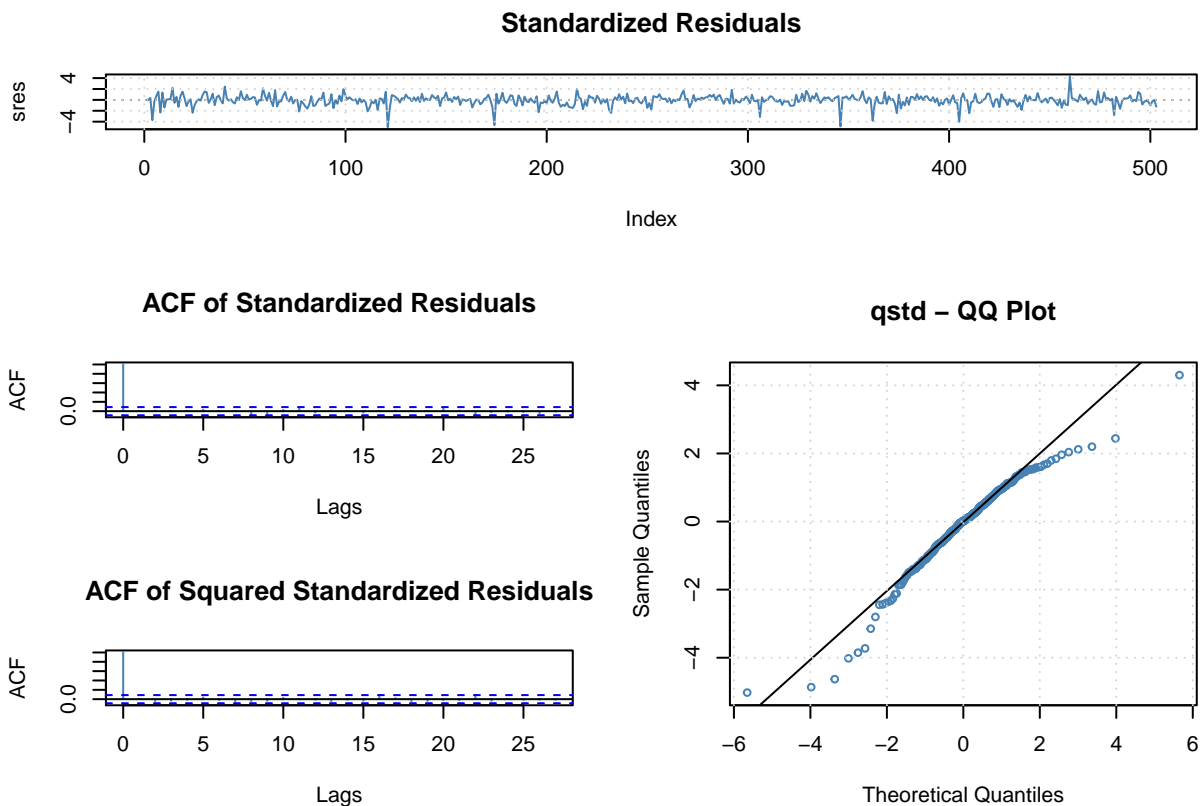
```
## Warning: Using formula(x) is deprecated when x is a character vector of length > 1.
## Consider formula(paste(x, collapse = " ")) instead.
```

```
summary(gf1)
```

```
##
## Title:
## GARCH Modelling
##
## Call:
## garchFit(formula = ~arma(2, 2) + garch(1, 1), data = qqqr, cond.dist = "std",
## trace = FALSE)
##
## Mean and Variance Equation:
## data ~ arma(2, 2) + garch(1, 1)
## <environment: 0x000000002028f0c8>
## [data = qqqr]
##
## Conditional Distribution:
## std
##
## Coefficient(s):
##      mu      ar1      ar2      ma1      ma2      omega
## 0.00013786 0.32645418 0.55550751 -0.44149180 -0.50703767 0.00000318
##      alpha1      beta1      shape
## 0.15315068 0.83699753 3.27412877
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      1.379e-04 7.813e-05 1.764 0.0777 .
## ar1      3.265e-01 3.373e-01 0.968 0.3332
## ar2      5.555e-01 2.868e-01 1.937 0.0528 .
## ma1     -4.415e-01 3.408e-01 -1.296 0.1951
## ma2     -5.070e-01 3.019e-01 -1.680 0.0930 .
## omega    3.180e-06 2.086e-06 1.525 0.1273
## alpha1   1.532e-01 6.236e-02 2.456 0.0140 *
## beta1    8.370e-01 5.450e-02 15.357 < 2e-16 ***
## shape    3.274e+00 5.515e-01 5.937 2.9e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## 1772.294 normalized: 3.523447
##
## Description:
## Tue May 25 19:42:05 2021 by user: orovi
##
## Standardised Residuals Tests:
##                               Statistic p-Value
## Jarque-Bera Test      R      Chi^2 645.5674 0
```

```
## Shapiro-Wilk Test R W 0.9224773 2.011307e-15
## Ljung-Box Test R Q(10) 5.670213 0.8421679
## Ljung-Box Test R Q(15) 10.91722 0.7584403
## Ljung-Box Test R Q(20) 13.63381 0.848567
## Ljung-Box Test R^2 Q(10) 5.907832 0.8229464
## Ljung-Box Test R^2 Q(15) 7.748837 0.9334568
## Ljung-Box Test R^2 Q(20) 14.57169 0.8003688
## LM Arch Test R TR^2 6.279338 0.901351
##
## Information Criterion Statistics:
## AIC BIC SIC HQIC
## -7.011109 -6.935592 -7.011735 -6.981484
```

```
layout(matrix(c(1,1,1,1,
                1,1,1,1,
                2,2,4,4,
                2,2,4,4,
                3,3,4,4,
                3,3,4,4),nrow=6, byrow=TRUE))
plot(gf1, which=9)
plot(gf1, which=10)
plot(gf1, which=11)
plot(gf1, which=13)
```



predict 14 days of future returns. recall this is a daily returns series, so these are predictions of future daily

returns.

```
### 14 day predictions of returns - PLOT  
preds <- predict(gf1, n.ahead=14, plot=TRUE)
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```


```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

```
## Warning in a_vec[i] <- ar[1:min(u2, i - 1)] * a_vec[(i - 1):(i - u2)] + : number  
## of items to replace is not a multiple of replacement length
```

Prediction with confidence intervals



preds

##	meanForecast	meanError	standardDeviation	lowerInterval	upperInterval
## 1	0.002305051	0.006739107	0.006739107	-0.01046200	0.01507210
## 2	0.001392226	0.006982061	0.006938888	-0.01183509	0.01461954
## 3	0.001872829	0.007176102	0.007131188	-0.01172209	0.01546775
## 4	0.001522642	0.007362891	0.007316616	-0.01242615	0.01547143
## 5	0.001675301	0.007543253	0.007495699	-0.01261518	0.01596578
## 6	0.001530606	0.007717683	0.007668898	-0.01309033	0.01615154
## 7	0.001568173	0.007886595	0.007836621	-0.01337276	0.01650910
## 8	0.001500058	0.008050350	0.007999227	-0.01375110	0.01675122
## 9	0.001498690	0.008209275	0.008157038	-0.01405355	0.01705093
## 10	0.001460404	0.008363659	0.008310342	-0.01438431	0.01730512
## 11	0.001447146	0.008513765	0.008459398	-0.01468194	0.01757623
## 12	0.001421550	0.008659828	0.008604443	-0.01498425	0.01782735
## 13	0.001405829	0.008802065	0.008745689	-0.01526943	0.01808109
## 14	0.001386479	0.008940671	0.008883331	-0.01555137	0.01832433

to produce pricing predictions the returns series needs to be applied to the tail of the price series, and then undifferenced and `exp()` applied.

```
# convert forecast means and intervals to prices
mf <- diffinv(preds$meanForecast, xi=log(qqq[503]))
```




```

mf <- mf[2:15] # drop first seed value (from original series)
li <- mf * (1 + preds$lowerInterval)
ui <- mf * (1 + preds$upperInterval)
mf <- exp(mf)
li <- exp(li)
ui <- exp(ui)

# get original time series plus the actual values for prediction interval
# create prediction, upper and lower interval series aligned with actuals
qqq_true <- window(qqqall, start='2017-07-01',end='2018-01-22')
qqq_mf <- tail(qqq_true, 14)
qqq_mf[,1] <- mf
qqq_li <- tail(qqq_true, 14)
qqq_li[,1] <- li
qqq_ui <- tail(qqq_true, 14)
qqq_ui[,1] <- ui

plt <- plot(ylim=c(130,175), qqq_true, type='l', main='QQQ predictions - ARMA(1,1)+GARCH(5,3)', ylab='Price')
plt <- lines(qqq_mf, lty=2, lwd=1, col='blue')
plt <- lines(qqq_li, lty=4, lwd=1, col='red')
plt <- lines(qqq_ui, lty=4, lwd=1, col='red')
plt <- addLegend('topleft', legend.names=c('predicted', 'lower', 'upper'),
               lty=c(2,4,4), lwd=c(1,1,1), col=c('blue', 'red', 'red'))
plt

```



compute RMSE of predictions

```
library(Metrics)
```

```
## Warning: package 'Metrics' was built under R version 4.0.5
```

```
##
```

```
## Attaching package: 'Metrics'
```

```
## The following object is masked from 'package:forecast':
```

```
##
```

```
## accuracy
```

```
rmse(tail(qqq_true,14), tail(qqq_mf,14))
```

```
## [1] 5.678657
```